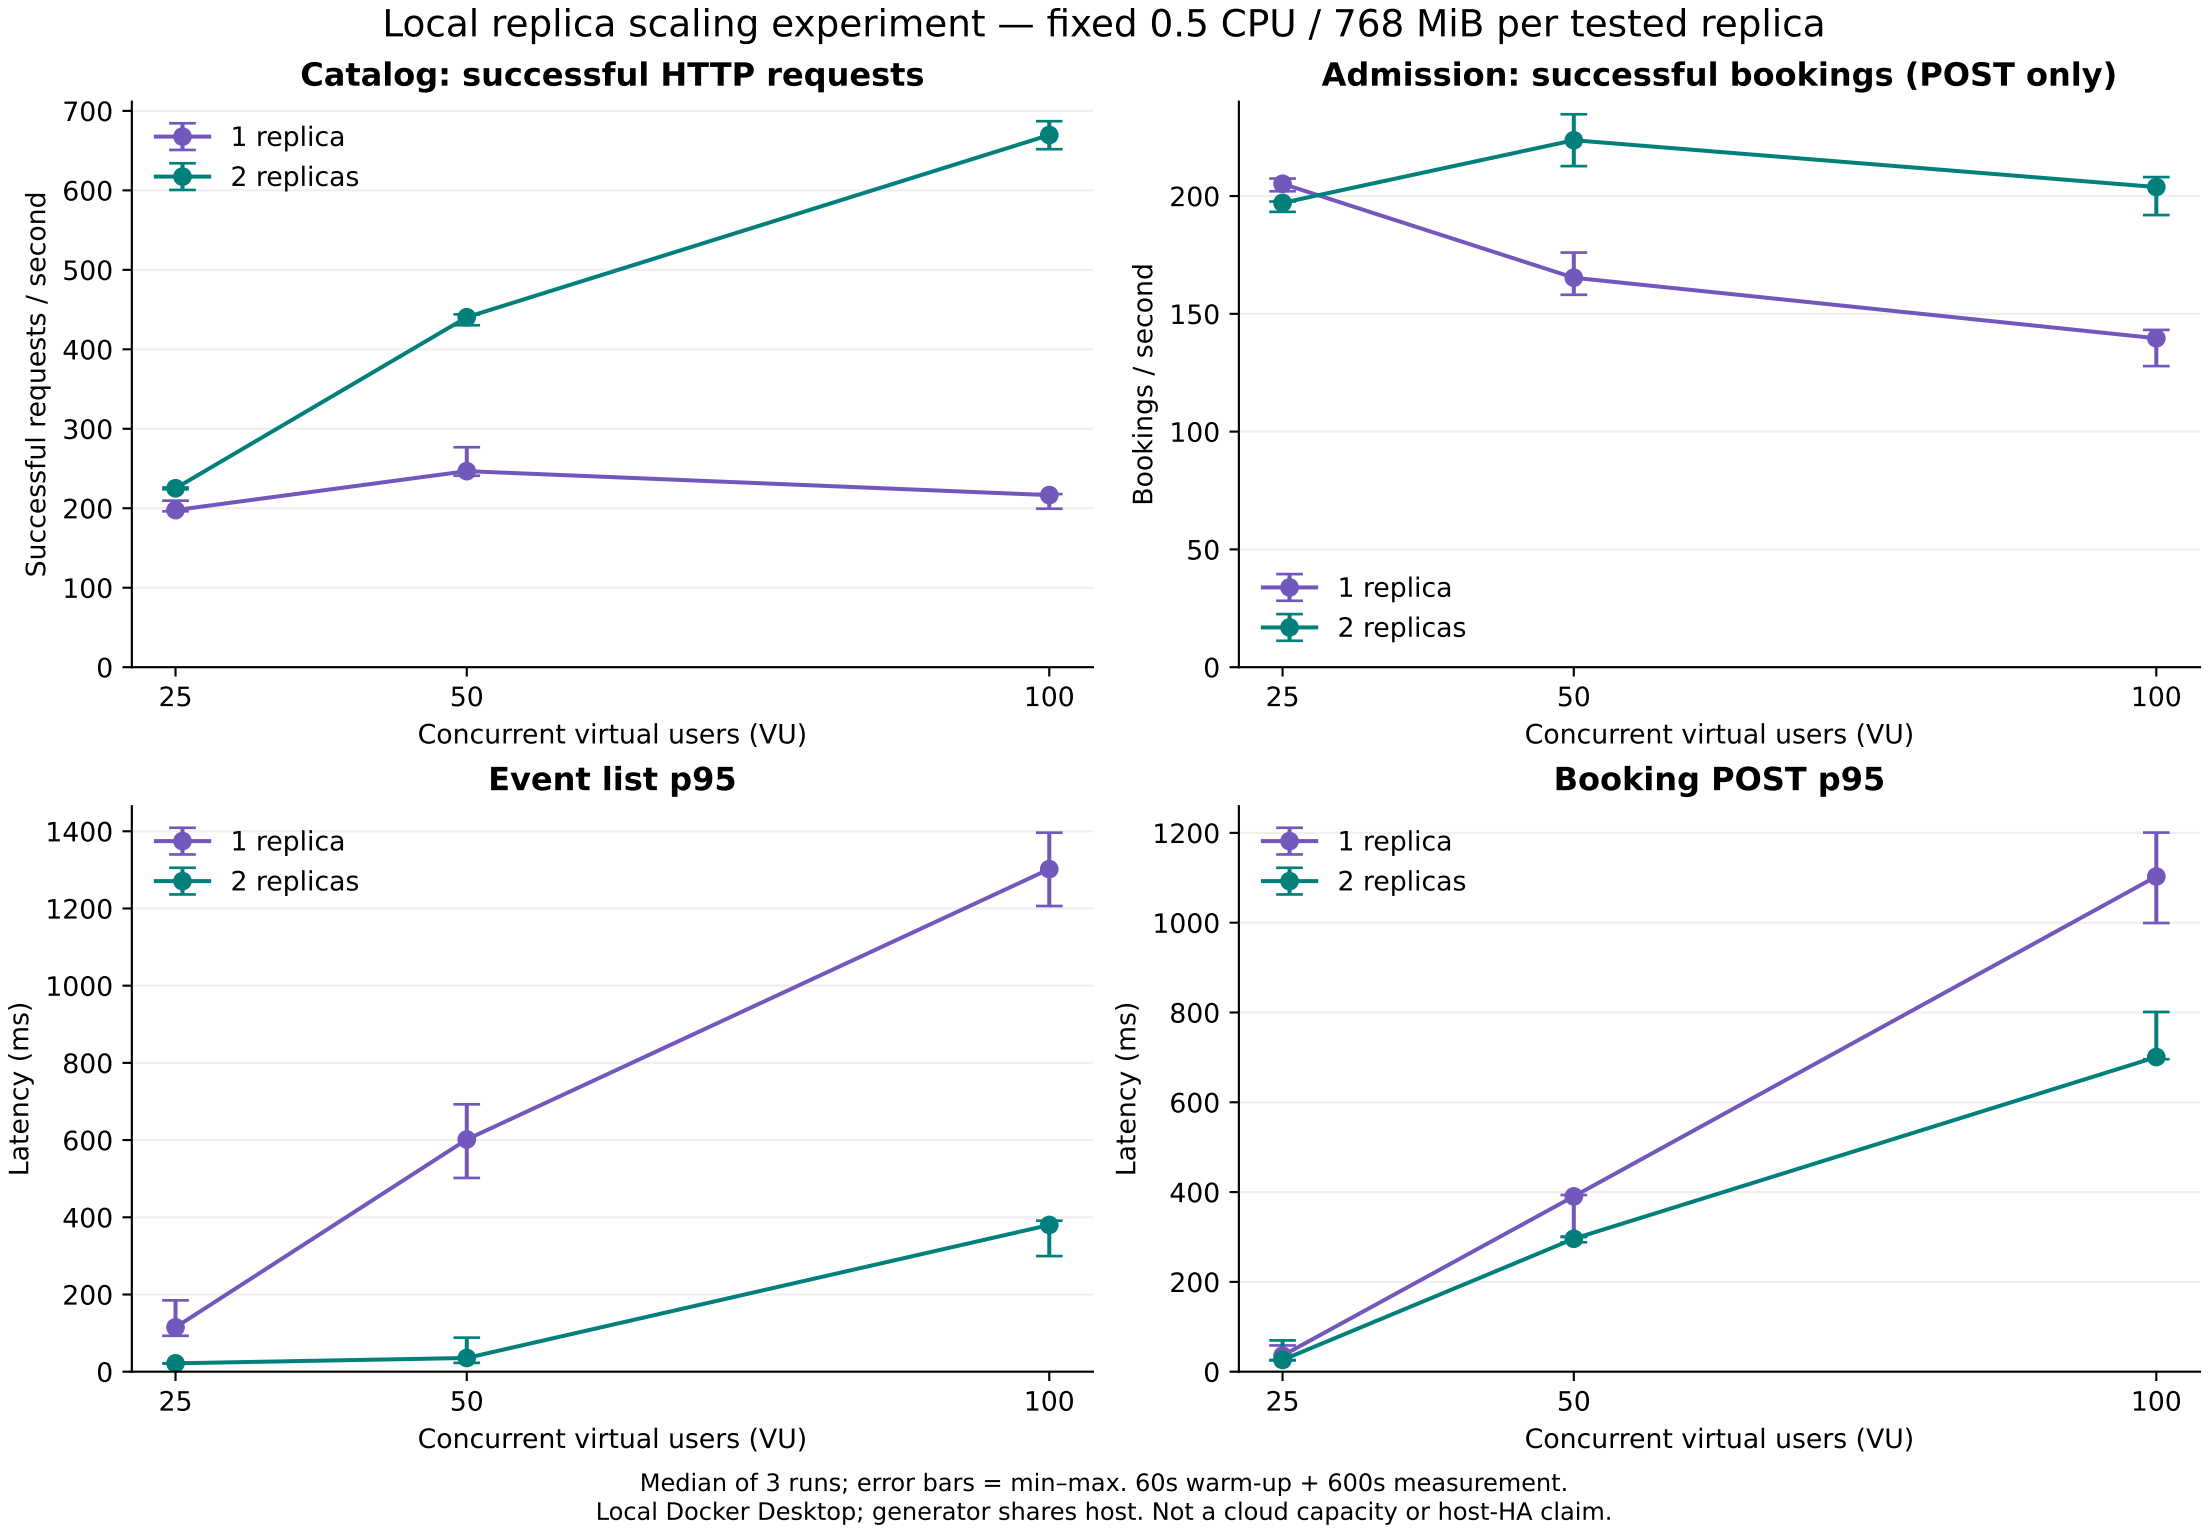
Local replica scaling experiment — fixed 0.5 CPU / 768 MiB per tested replica
Catalog: successful HTTP requests
Successful requests / second
Concurrent virtual users (VU)
1 replica
2 replicas
Admission: successful bookings (POST only)
Bookings / second
Concurrent virtual users (VU)
Event list p95
Latency (ms)
Concurrent virtual users (VU)
Booking POST p95
Latency (ms)
Concurrent virtual users (VU)
Median of 3 runs; error bars = min–max. 60s warm-up + 600s measurement.
Local Docker Desktop; generator shares host. Not a cloud capacity or host-HA claim.